Migrate PostgreSQL Single Server to Flexible Server using Azure portal

Article • 01/12/2022 • 8 minutes to read •

In this article

This article shows you how to create a migration from your Azure database for PostgreSQL single server to flexible server using our automated migration service through Azure portal.

🔗 Pre-requisites

Make sure to take care of the pre-requisites listed here, which are necessary to get started with the automated migration service.
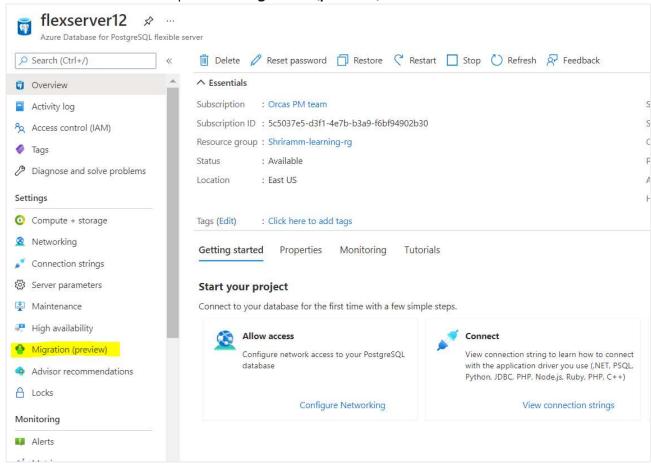
Portal Experience

The automated migration service comes with a simple, wizard-based portal experience to create a migration from single server to flexible server.
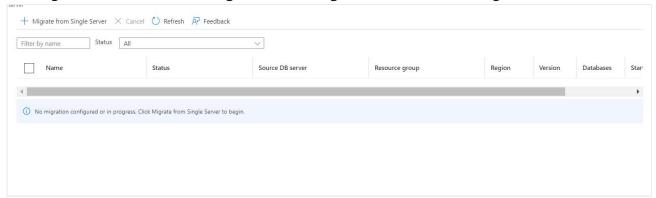
Sign into the Azure portal

Open your web browser and go to the portal . Enter your credentials to sign in to the portal. The default view is your service dashboard. If you don't have an Azure subscription, create a free Azure account

If you haven't created an Azure database for PostgreSQL flexible server, go ahead and create one using this link. Else, find and navigate to your flexible server instance.

Once you are in the **Overview** tab of your flexible server, use the left navigation window and scroll down to the option of **Migration (preview)** and click on it.

If this is the first time you are using the migration service, you will see an empty grid with a message to click the button **Migrate from Single Server** to start a migration.

If you have already created migrations to your flexible server, you should see the list of migrations that were attempted to this flexible server from single servers.

Click on the **Migrate from Single Server** option. You'll be taken through a wizard-based setup to create a migration to this flexible server from any single server.

Setup Tab

Migrate from single server into this PostgreSQL flexible server ···
Microsoft - preview

Setup   Source   Target   Networking   Review + create

Pre-requisites -
1) Create an Azure Active Directory application. How to do? ⤢
2) Create a new client secret for your azure active directory application. How to do? ⤢

We will be using a Migration Resource group. This is the resource group where all the migration related components will be created. By default it is resource group of the target flexible server and all the components will be cleaned up automatically once the migration completes.

3) Assign contributor roles to source server, target server, migration resource group. Note in case of private access for source/target server, add Contributor privileges to the corresponding VNet. How to do? ⤢

Restrictions -
1) Recommended to use only for individual DBs <= 1 TB.
2) All logical replication restrictions ⤢ in PostgreSQL apply.
3) You can migrate up to 8 databases from a server in a single activity. If you need to migrate more, create multiple activities.

Migration name * ⓘ	
Migration Resource Group * ⓘ	Shriramm-learning-rg ⌄
Azure Active Directory App * ⓘ	Select

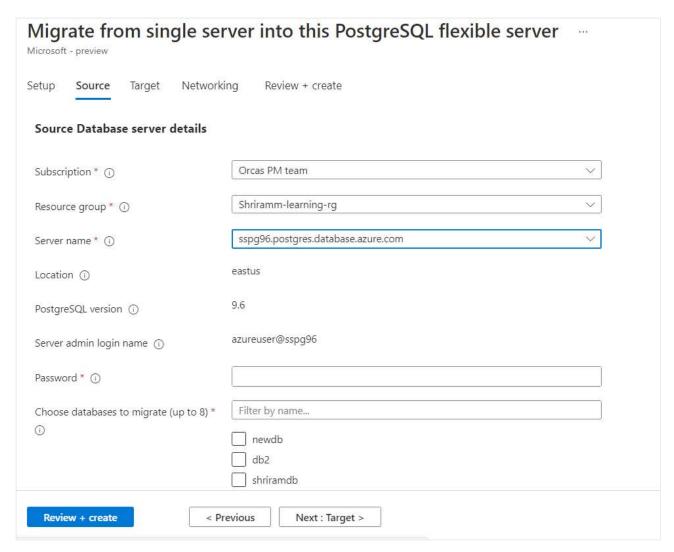[Review + create]   [Next : Source >]

There is a **pre-requisites** section, the details of which are documented in detail here There is also a **restriction** section that lists down the restrictions that are applicable to this migration service

- The **Migration name** field accepts only alphanumeric characters and does not accept any special characters except '-'. The name can't start with a '-' and should be unique for a target server. No 2 migrations to the same flexible server can have the same name.
- The **Migration resource group** is where all the migration-related components will be created by the migration service. By default it's resource group of the target flexible server and all the components will be cleaned up automatically once the migration completes. If you want to create a temporary resource group for migration-related purposes, create a resource group and select the same from the dropdown.
- For the **Azure Active Directory App**, click the **select** option and pick the app that was created as a part of the pre-requisite step. Once the AAD App is chosen, paste the client secret that was generated for the AAD app to the **Azure Active Directory Client Secret** field.

Migrate from single server into this PostgreSQL flexible server ···

Microsoft - preview

Setup Source Target Networking Review + create

Pre-requisites -
1) Create an Azure Active Directory application. How to do? 🗗
2) Create a new client secret for your azure active directory application. How to do? 🗗

We will be using a Migration Resource group. This is the resource group where all the migration related components will be created. By default it is resource group of the target flexible server and all the components will be cleaned up automatically once the migration completes.

3) Assign contributor roles to source server, target server, migration resource group. Note in case of private access for source/target server, add Contributor privileges to the corresponding VNet. How to do? 🗗

Restrictions -
1) Recommended to use only for individual DBs <= 1 TB.
2) All logical replication restrictions 🗗 in PostgreSQL apply.
3) You can migrate up to 8 databases from a server in a single activity. If you need to migrate more, create multiple activities.

Migration name * ⓘ	mymigraton1 ✓
Migration Resource Group * ⓘ	Shriramm-learning-rg ⌄
Azure Active Directory App * ⓘ	aad-migration-demo ✎
Azure Active Directory Client Secret * ⓘ	•••••••••••••••••••••••••• ✓ ✅ This field is required.

[Review + create]   [Next : Source >]

Click on the **Next** button.

Source tab

Migrate from single server into this PostgreSQL flexible server ···
Microsoft - preview

Setup   **Source**   Target   Networking   Review + create

Source Database server details

Subscription * ⓘ	Orcas PM team ⌄
Resource group * ⓘ	Shriramm-learning-rg ⌄
Server name * ⓘ	sspg96.postgres.database.azure.com ⌄
Location ⓘ	eastus
PostgreSQL version ⓘ	9.6
Server admin login name ⓘ	azureuser@sspg96
Password * ⓘ	
Choose databases to migrate (up to 8) * ⓘ	Filter by name...

☐ newdb
☐ db2
☐ shriramdb

Review + create   < Previous   Next : Target >

In this tab, you'll be asked to give details related to the source single server. As soon as you pick the **Subscription** and **Resource Group**, the dropdown for server names will have the list of single servers under that particular resource group. Choose any single server from the drop down.

Once the single server is chosen, the fields such as **Location**, **PostgreSQL version**, **Server admin login name** are automatically pre-populated. The server admin login name is the admin username that was used to create the single server. You should also see the list of databases inside the single server populated as shown in the above pic.

Enter the corresponding password for the admin user and pick the list of databases that you want to migrate. A max of eight databases can be chosen per migration. Once the list of databases is selected, click on the **Next** Button

Target tab

This tab displays metadata of the flexible server like the **Subscription**, **Resource Group**, **Server name**, **Location**, and **PostgreSQL version**. It displays **server admin login name** which is the username that was used during the creation of the flexible server.

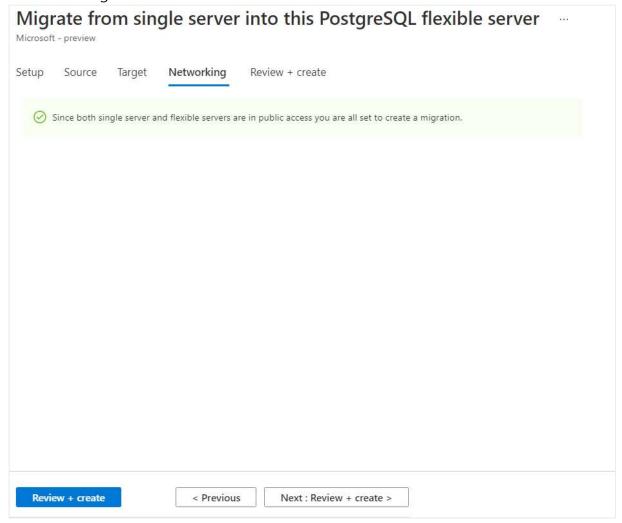Enter the corresponding password for the admin user.

Choose an option **yes/no** for **Authorize DB overwrite**. If you set the option to **Yes**, you give this migration service permission to overwrite existing data in case when a database that is being migrated to flexible server is already present. If set to **No**, it goes into a waiting state and asks you for permission either to overwrite the data or to cancel the migration.

Click on the **Next** button

Networking tab

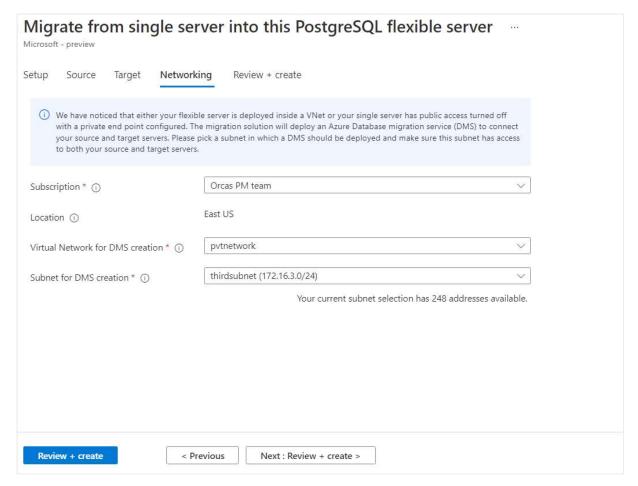The content on the Networking tab depends on the networking topology of your source and target servers.

- If both source and target servers are in public access, then you are going to see the below message

In this case, you need not do anything and can just click on **Next** button.

- If either the source or target server is in private access, then the content of the networking tab is going to be different. Let us try to understand what does private access mean for single server and flexible server
  - **Single Server Private Access - Deny public network access** set to **Yes** and a private end point configured
  - **Flexible Server Private Access -** When flexible server is deployed inside a VNet.

  If either source or target is private access, then the networking tab looks like the following

Migrate from single server into this PostgreSQL flexible server ...
Microsoft - preview

Setup   Source   Target   **Networking**   Review + create

ⓘ We have noticed that either your flexible server is deployed inside a VNet or your single server has public access turned off with a private end point configured. The migration solution will deploy an Azure Database migration service (DMS) to connect your source and target servers. Please pick a subnet in which a DMS should be deployed and make sure this subnet has access to both your source and target servers.

Subscription * ⓘ   Orcas PM team

Location ⓘ   East US

Virtual Network for DMS creation * ⓘ   pvtnetwork

Subnet for DMS creation * ⓘ   thirdsubnet (172.16.3.0/24)

Your current subnet selection has 248 addresses available.

[Review + create]   [< Previous]   [Next : Review + create >]

All the fields will be automatically pre-populated with subnet details. This is the subnet in which the migration service will deploy Azure DMS to move data between the source and target.

- If source has private access and target has public access, then the pre-populated subnet details would be that of the same subnet for which the private end point is configured.
- If source has public access and target has private access, then the pre-populated subnet details would be that of the subnet inside the same VNet inside which the flexible server is deployed. It will not be the same subnet as that of flexible server since flexible server needs delegated subnets
- If both source and target are in private access, then the pre-populated subnet details would be that of the subnet inside the same VNet inside which the flexible server is deployed.
  You can go ahead with the suggested subnet or choose a different subnet. But make sure that the selected subnet can connect to both the source and target servers.
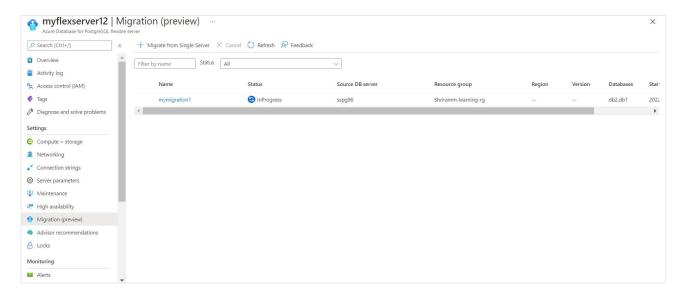
After picking a subnet, click on **Next** button

Review + Create tab

This tab gives a summary of all the details given by the user for creating the migration. Review the details and click on the **Create** button.

After clicking on the **Create** button, you should see a notification saying the migration was successfully created in a few seconds
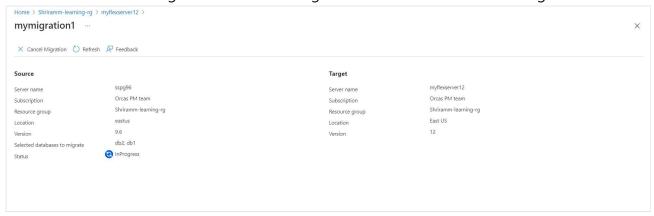

You should be automatically redirected to **Migrations Preview** page of flexible server that will have a new entry of the recently created migration

The grid displaying the migrations has various columns including **Name**, **Status**, **Source server name**, **Region**, **Version**, **Database names**, and the **Migration start time**. By default, the grid shows the list of migrations in the decreasing order of migration start time. In other words, the more recently created migrations appear on top of the grid.

You can use the refresh button to refresh the status of the migrations.

You can click on the migration name in the grid to see the details of that migration.
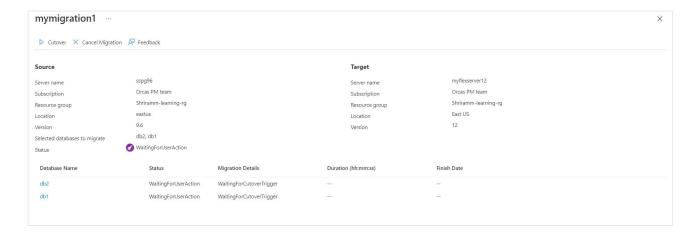

As soon as a migration is created, it moves into the state of **InProgress**. During this state, a bunch of pre-requisites steps will be carried out by the migration service. This includes creating an Azure DMS, ensuring DMS has access to both servers by adding its IP to the list of firewall rules and few other maintenance tasks. In general, it takes around 8 to 10 minutes for the migration to move from the **InProgress** state to **Migrating Data** state.
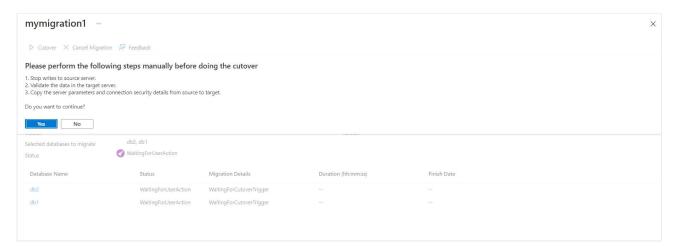
Migrating Data is the state where the dump and restore of schema happens followed by movement of data from source to target server. Once this is completed, the migration moves to the state of **WaitingForUserAction** and substate of **WaitingForCutOver**.

You can click on the migration name to go into the migration details page and should see the substate of **WaitingForCutover**.
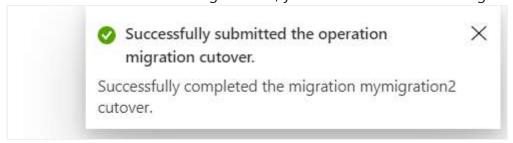
The migration moves into the substate **WaitingForCutOver** once it has finished the dump and restore of the databases of your single server. At this stage, the ongoing writes at your source server will be replicated to the target flexible server using the logical replication feature of PostgreSQL. Click on the **Cutover** button and you should see the following message

Make sure to stop writes on your server and click on the **Yes** button to start cutover.

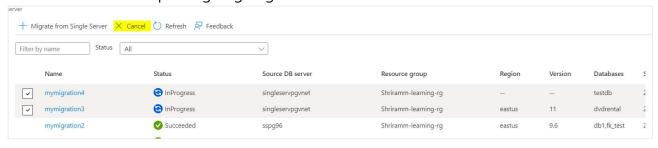In a few seconds after starting cutover, you should see the following notification

Once the cutover is complete, the migration moves to **Succeeded** state and migration of schema data from your single server to flexible server is now complete. You can use the refresh button in the page to check if the cutover was successful.

After verifying the data in your target server, you can make changes to your application to point to the flexible server.

You also have the option to cancel any ongoing migrations. For a migration to be canceled, it has to be in **InProgress** or **WaitingForUserAction** state. You can't cancel a migration that has either **Succeeded** or **Failed**.

You can choose multiple ongoing migrations at once and can cancel them.

Note that **cancel migration** just stops any more further migration activity on your target server. It will not drop or roll back any changes on your target server that were done by the migration attempts.

Post Migration

Once your migration completes, follow the post migration steps listed here

Next steps

- Understand what's available for compute and storage options - Flexible server
- Learn about supported PostgreSQL Database Versions in Flexible Server
- Learn about current limitations in Flexible Server